MEMORANDUM

CREDRESOLVE

TO: Leadership Team
FROM: Nirav Gupta, Collections Analytics Review
DATE: September 2, 2026
RE: Verification of the Reported 11% Recovery Improvement

What happened

The reported figure of **₹134.1 Cr** in monthly recovery is wrong. The corrected figure is ₹107.5 Cr a gap of [₹26.7 Cr, or 24.8%]. Most of that gap comes from payment retries: 1,931 successful payments shared a transaction reference with an earlier payment already counted, and a further 486 rows were byte-identical duplicates sitting in the source system. Once those are removed and ₹9.5 Cr in reversed payments is netted out, the true number falls well short of the reported one.

The "11% month-on-month improvement" does not survive the correction either. It was never a trend. It was one comparison February to March pulled out of a series that averaged -0.4% month-on-month across the period and moved **-4.5%** from January to July. Two adjacent months in a noisy series will occasionally show a double-digit swing; that swing became the headline number, and the months on either side of it were left out of the sentence.

A second reported figure collapses under the same scrutiny. Promise-to-pay records list 4,489 promises as "kept." Checking each one against an actual payment within seven days of the promised date leaves 567 -12.6% of the labelled count. The other 3,922 were marked kept on the strength of the label alone, with no payment behind them.

Why it happened

Three separate mechanisms are doing the work here, and they are not the same mechanism.

The first is a straightforward data defect: duplicate payment ingestion. This alone accounts for the ₹26.7 Cr gap and does not require any story about borrower behavior or agent performance to explain.

The second is a reporting artifact, not a data defect: the 11% figure was a cherry-picked pair of months presented as a trend. Nothing was manipulated in the underlying records to produce this one it is a selection problem in how the number was chosen for the report.

The third is more interesting and harder to dismiss. The aggregate recovery-per-paying-account looks flat (+[0.4% January to July]), but that stability is hiding two segments moving in opposite directions. Low-risk accounts show recovery-per-account rising 6.4% over the period; high-risk accounts are down 4.9%. Older-vintage loans (opened before 2025) are recovering better than 2025-vintage loans. Neither shift is large enough to move the aggregate, but together they describe a portfolio quietly getting harder to collect on newer business while the easier back-book compensates. That is a real underlying pattern, not a data artifact, and it is worth more attention than the fake 11% ever received.

Two further tests are worth stating because they rule things out rather than in. Portfolio composition by risk segment barely moved (5.0 percentage points, max) across the period, so the reported improvement cannot be attributed to a change in the mix of accounts being worked. The collectible population - accounts still active, delinquent, on a promise, or non-performing, tracked point-in-time rather than by current status held at 56.7–57.4% every month. Written-off accounts were not quietly leaving the denominator to inflate a conversion rate. Both of those are candidate explanations that a skeptical reader might reach for first, and both check out clean.

How confident we are

Graded by evidence standard, not uniformly: the duplicate-payment inflation, the flat-to-declining trend, and the denominator stability are established directly from the transaction data and carry no material uncertainty - call them fact. The segment divergence (**Simpson's paradox pattern**) and the cohort effect are correlational: real and consistent in the data, but the underlying risk-segment and DPD fields in this dataset show signs of partial randomization at the account level, so a causal story about *why* newer-vintage accounts are underperforming should wait for cleaner covariates.

One finding needs a caveat attached before it goes further: comparing targeted accounts (40.2% pay rate) against never-targeted accounts (**41.4%**) suggests the current targeting strategy adds no measurable lift, or possibly a small negative one. That comparison is unconditional it does not yet control for whether the never-targeted group differs systematically in risk or DPD composition from the targeted group. The result deserves to be treated seriously, but not yet as settled.

A second limitation belongs here rather than buried in a footnote: the **12.6%** "verified kept" figure for promises-to-pay depends on a seven-day window applied in both directions from the promised date. Widening or narrowing that window moves the figure from **0.2% to 6.1%** across a 1-to-30-day sweep. The number is real, but it is also a choice, and a different analyst using a different window would report a different percentage from the same records.

What we should do

Two actions, and they are not the same kind of action. First, stop reporting **₹134.1 Cr** and **+11%**. Replace them with **₹107.5 Cr** and a flat-to-declining trend in every dashboard and leadership update currently citing the old figures this is a correction to make immediately, independent of anything else in this memo.

Second, hold the **₹10 Cr** allocation decision. Of the six candidate investments leadership is weighing, only one borrower targeting has been tested against this data, and the test result argues against investing further in the current targeting approach without first understanding why it shows no lift. The other five channels have not been evaluated at all. Committing capital now would mean choosing among six options having genuinely examined one. The next concrete step is a controlled before/after comparison around the point where the targeting strategy changed mid-year, using the untargeted accounts as a natural control group that analysis is not yet complete, and it should be, before **₹10 Cr** moves anywhere.

Financial impact

The immediate, quantified impact of this review is the **₹26.7 Cr** correction to the reported recovery base money that does not exist and should not appear in any forward projection built on the old number. Beyond that correction, a defensible expected-return figure for the **₹10 Cr** allocation cannot be produced yet. Doing so now, on five untested channels and one channel showing a flat-to-negative signal, would mean presenting a confident number without confident evidence behind it precisely the failure mode this review was commissioned to catch.